



ELEMENTS OF DATA SCIENCE AND STATISTICAL LEARNING

SPRING 2017

Week 8

OUTLINE


Unsupervised learning, continued

- How many clusters?
- Measures of “goodness of clustering”
- Examples

TEST DATASET

- Let us consider again our artificial example. We will use it to develop some intuition (and code)

```
x=c(rnorm(30,mean=2),rnorm(50,mean=-2),rnorm(40,mean=4))  
y=c(rnorm(30,mean=-2),rnorm(50,mean=2),rnorm(40,mean=2))  
kf=kmeans(cbind(x,y),3)  
plot(x,y,pch=19,col=kf$cluster)  
l.cl=c(rep("A",30),rep("B",50),rep("C",40))
```




WITHIN-CLUSTER VARIATION

- Within-cluster variation $W = \sum_{k=1}^K \sum_{C(i)=k} \|x_i - \bar{x}_k\|_2^2 \leftarrow (\text{sum of squares}): \mathbf{v} = (v_1, \dots, v_p) \Rightarrow \|\mathbf{v}\|_2 = \sqrt{\sum_{j=1}^p v_j^2}$
- Where the cluster centers are defined as $\bar{x}_k = \frac{1}{n_k} \sum_{C(i)=k} x_i$
- Data points x_i are considered here to be multi-dimensional, i.e. each point = a vector of observations for p variables


```
w=numeric(20)
for ( k in 1:20 ) {
  kf=kmeans(cbind(x,y),k)
  w[k] = kf$tot.withinss
}
plot(1:20,w,type="b",lwd=2,pch=19,xlab="K",
      ylab=expression("SS[within]"))
```

elbow




WITHIN CLUSTER VARIATION ON RANDOM DATA

- How striking is the elbow?
- Striking compared to what?
- For example, results of the same algorithm (K-means clustering) applied to data that is *known* to have *no clusters*
 - Random sample from normal or uniform distribution
- Of course, distribution we sample has to be roughly representative of the data we work with
 - E.g. standard normal has much lower variance than our simulated dataset
- Random sample of uniform distribution bounded by range of every attribute is commonly used alternative
- In this case total within cluster variance decreases faster for original data as compared to random sample




ORIGINAL, NORMAL AND UNIFORM EXAMPLES


Original



Normal



Uniform



- Sample from bivariate standard normal is obviously much tighter than that in the “original” data
- Uniform sample is more representative of “the same data without clusters”

CODE TO GENERATE PREVIOUS PLOTS

```
plot(1:20,w,type="b",lwd=2,pch=19,xlab="K",
     ylab=expression("SS[within]"))
for ( i in 1:10 ) {
  wrnd = numeric()
  for ( k in 1:20 ) {
    krnd =
kmeans(matrix(rnorm(2*length(x)),ncol=2),k)
    wrnd[k] = krnd$tot.withinss
  }
  points(wrnd,type="l",col="blue")
}
```

```
for ( i in 1:10 ) {
  wrnd = numeric()
  for ( k in 1:20 ) {
    krnd =
kmeans(apply(cbind(x,y),2,function(x)runif(length(x),min(x),max(x))),k)
    wrnd[k] = krnd$tot.withinss
  }
  points(wrnd,type="l",col="red")
}


legend("topright",c("ori","norm","unif"),col=c(
"black","blue","red"),
      text.col=c("black","blue","red"),pch
h=c(19,NA,NA),lty=1,lwd=c(2,1,1))
```

```
old.par <- par(mfrow=c(1,3),ps=16)
plot(x,y,pch=19,col=kf$cluster,main="Original")
plot(matrix(rnorm(2*length(x)),ncol=2),xlim=range(x),ylim=range(y),main="Normal")
plot(apply(cbind(x,y),2,function(x)runif(length(x),min(x),max(x))),main="Uniform")
par(old.par)
```

BETWEEN-CLUSTER VARIATION


- Between-cluster variation, $B = \sum_{k=1}^K n_k \|\bar{x}_k - \bar{x}\|_2^2$
- Note that the *total sum of squares* $T = \sum_{i=1}^N \|x_i - \bar{x}\|_2^2 = B + W$

```
w=numeric(20)
for ( k in 1:20 ) {
  kf=kmeans(cbind(x,y),k)
  w[k] = kf$betweenss
}
plot(1:20,w,type="b",lwd=2,pch=19,xlab="K",
     ylab=expression("SS[between]"))
```



BETWEEN-CLUSTER VARIATION ON RANDOM DATA

- We know now that standard normal is poor comparator for the data with three clusters
- For samples from uniform distribution between cluster sum of squares also increases rapidly for the first few clusters
- And even to higher levels on average than what is observed for higher values of K



CODE FOR THE PREVIOUS PLOT

```
plot(1:20,w,type="b",lwd=2,pch=19,xlab="K",
     ylab=expression("SS[between]"),ylim=range(w)*c(1,1.2))

for ( i in 1:10 ) {
  btwrnd = numeric()
  for ( k in 1:20 ) {
    krnd =
kmeans(matrix(rnorm(2*length(x)),ncol=2),k)
    btwrnd[k] = krnd$betweenss
  }
  points(btwrnd,type="l",col="blue")
}
```

```
for ( i in 1:10 ) {
  btwrnd = numeric()
  for ( k in 1:20 ) {
    krnd =
kmeans(apply(cbind(x,y),2,function(x) runif(
length(x),min(x),max(x))),k)
    btwrnd[k] = krnd$betweenss
  }
  points(btwrnd,type="l",col="red")
}

legend("topleft",c("ori","norm","unif"),col
=c("black","blue","red"),
      text.col=c("black","blue","red"
),pch=c(19,NA,NA),lty=1,lwd=c(2,1,1))
```

CH INDEX

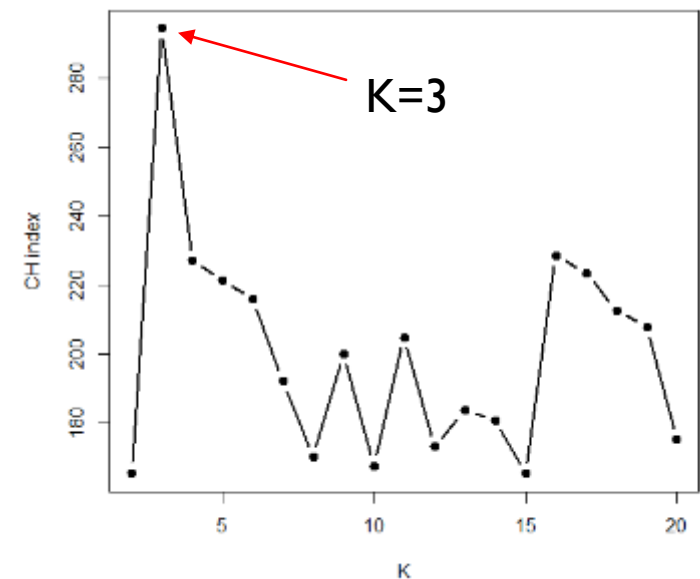
- Consider a score that attempts to strike a balance between within- and between-cluster sums of squares:

$$CH(K) = \frac{B(K)/(K-1)}{W(K)/(n-K)}$$

compare this to *F*-statistic (e.g. for linear regression it's Eq(3.23) in ISLR): $F = \frac{(TSS-RSS)/p}{RSS/(n-p-1)}$

- Called CH index [Calinski & Harabasz (1974), "A dendrite method for cluster analysis"]
- Calculate in a (reasonable) range of K values, pick K that maximizes the index:


```
w=numeric(20)
for ( k in 2:20 ) {
  kf=kmeans(cbind(x,y),k)
  w[k] = (kf$betweenss/(k-1)) /
        (kf$tot.withinss/(length(x)-k))
}
plot(2:20,w[-1],type="b", lwd=2,pch=19,xlab="K",
     ylab="CH index")
```



CH-INDEX ON RANDOM DATA

- CH-index achieves much higher values on original as compared to random data from uniform distribution with the same ranges
- Tends to gradually increase with the increase of K
- Value of K corresponding to the maximum in CH-index does not imply that this is the number of “real” clusters

```
for ( i in 1:10 ) {  
  chrnd = numeric()  
  for ( k in 1:20 ) {  
    krnd =  
kmeans(apply(cbind(x,y),2,function(x) runif(len  
gth(x),min(x),max(x))),k)  
    chrnd[k] = (krnd$betweenss/(k-1))/  
      (krnd$tot.withinss/(length(x)-k))  
  }  
  points(chrnd,type="l",col="red")  
}
```



GAP STATISTICS

- Idea: compare against the “null distribution”
 - What is “null”?
 - No clusters = the points are distributed randomly and uniformly over an encapsulating box
 - Let’s try it. Are those real clusters? (remember, this is how random points “cluster”!)

```
m.new=apply(cbind(x,y),2,function(x) {  
  runif(length(x),min=min(x),max=max(x))  
})
```

```
m.new[1:3,]
```

	x	y
[1,]	0.3242968	2.114528
[2,]	4.8677956	3.327992
[3,]	3.6180172	4.437447

```
kf=kmeans(m.new,3)
```

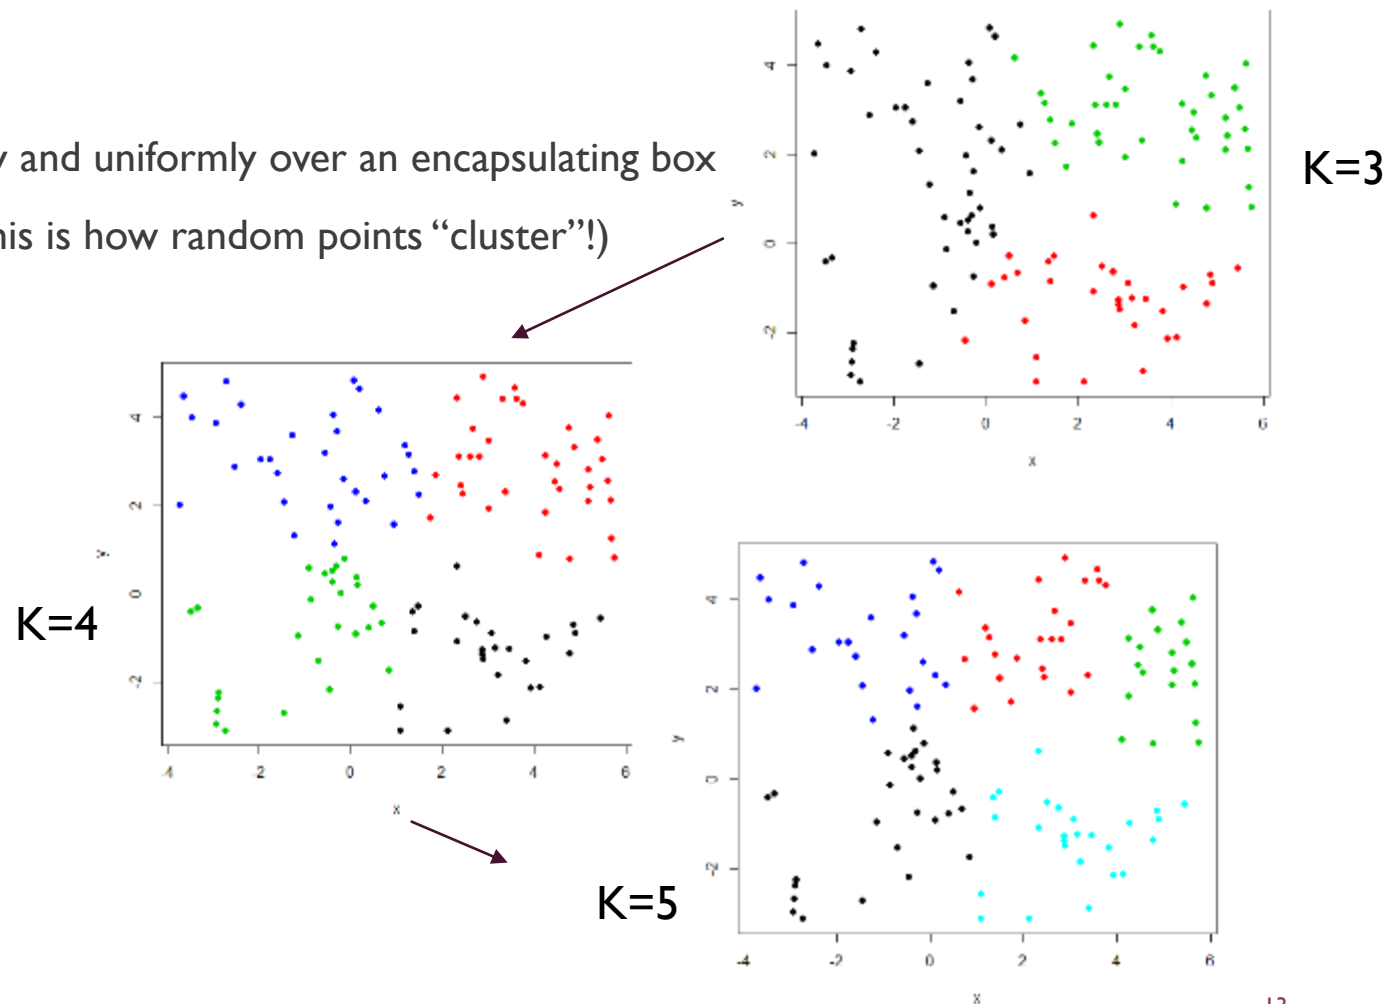
```
plot(m.new,pch=19,col=kf$cluster)
```

```
kf=kmeans(m.new,4)
```

```
plot(m.new,pch=19,col=kf$cluster)
```

```
kf=kmeans(m.new,5)
```

```
plot(m.new,pch=19,col=kf$cluster)
```



GAP STATISTICS: CONTINUED

- Let us use the ratio of the “null” within-cluster sum of squares $W_{\text{unif}}(K)$ (computed from uniformly scattered points) to the actual within-cluster SS, $W(K)$ (computed on real data) as the measure of “goodness of clustering”
 - With no cluster structure we expect the clusters to be “bad” and wide \rightarrow large W_{unif} (but still decreases with K !)
 - If there are true clusters in the data, we expect them to be “tighter” than those “seen” in uniform random data, so $W_{\text{unif}}(K)/W(K)$ is large.
 - Since we are computing a quantity (W_{unif}) from random data, it is a good idea to perform a few resamplings and take the average, so below we consider W_{unif} to be such average, with standard error $s(K)$
 - It is convenient to take log: $\text{Gap}(K) = \log W_{\text{unif}}(K) - \log W(K)$
 - We define the optimal K as the value where gap reaches its first maximum:
 - $K_{\text{opt}} = \min\{K \in \{1 \dots K_{\text{max}}\}: \text{Gap}(K) \geq \text{Gap}(K + 1) - s(K + 1)\}$


GAP STATISTICS: CODE

- Let's develop the code for gap statistics:

```
# takes matrix of observations of p variables (points in
# p-dimensional space: row=point), generates the
# same number of p-dimensional points randomly and
# uniformly scattered in the bounding box:
lw.unif=function(m,K,N=20) {
  w=numeric(N)
  for ( i in 1:N ) {
    m.new=apply(m,2,function(x) {
      runif(length(x),min=min(x),max=max(x))
    })
    kf=kmeans(m.new,K)
    w[i] = kf$tot.withinss
  }
  return( list(LW=mean(log(w)),SE=sd(log(w))/sqrt(N)) )
}
```


```
# computes the gap and the  $\text{LogW}_{\text{unif}}$  SE
# for different K:
gap=numeric(20)
se = numeric(20)
for ( k in 1:20 ) {
  kf=kmeans(cbind(x,y),k)
  sim = lw.unif(cbind(x,y),k)
  gap[k] = sim$LW - log(kf$tot.withinss)
  se[k] = sim$SE
}
plot(1:20,gap,pch=19,type="b")
# find optimal K:
min(which(gap[-length(gap)] >=
  (gap-se)[-1]))
[1] 3
```

Note: gap statistics is implemented in packages lga and SAGx



ORIGINAL AND NULL WITHIN-CLUSTER SUM OF SQUARES

```
se = numeric(20)
rndlw = numeric(20)
orilw = numeric(20)
for ( k in 1:20 ) {
  kf=kmeans(cbind(x,y),k)
  sim = lw.unif(cbind(x,y),k)
  rndlw[k] = sim$LW
  orilw[k] = log(kf$tot.withinss)
  se[k] = sim$SE
}
plot(1:20,rndlw,type="l",lwd=1,col="red",xlab="K",
     ylab="log(withinss)")
arrows(1:20, rndlw-se, 1:20, rndlw+se,
       length=0.05, angle=90, code=3, col=2)
points(orilw,type="b",pch=19,lwd=2)
legend("topright",c("ori","unif"),text.col=1:2,col=1:2,
      pch=c(19,3),lty=1,lwd=c(2,1))
```



- According to gap statistics, the “best” number of clusters is where the gap between original and null is the largest


SILHOUETTE

- Consider a point x_i . *Silhouette* is defined for each individual point in the following way:
 - Let $a(i)$ be the average dissimilarity (i.e. distance) between x_i and all other points of the cluster $C(i)$ to which x_i belongs
 - If x_i is in its own cluster (with no other points in it), then the silhouette is $s(i)=0$; otherwise
 - For each other cluster C (i.e. excluding $C(i)$) let $d(i,C)$ be the average distance from x_i to all observations in C
 - Define the “distance to the closest cluster” as $b(i)=\min_C d(i,C)$
 - Define silhouette as

$$s(i) = \frac{b(i) - a(i)}{\max(a(i), b(i))}$$

- Properties:


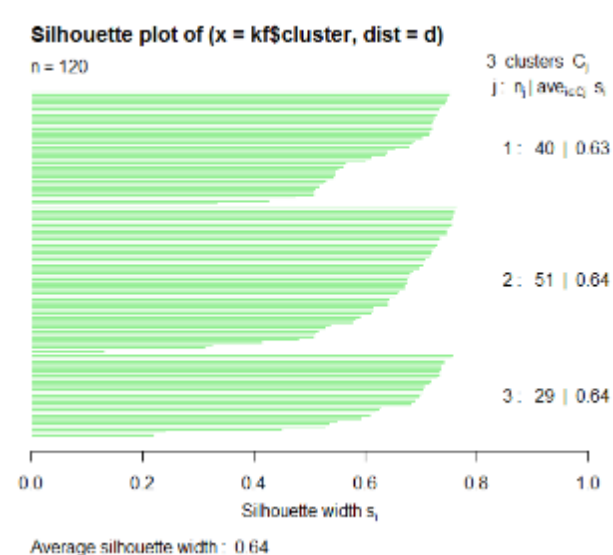
- Normalized!
- If “average width” $a(i)$ is small and the distance to the closest cluster $b(i)$ is large, then $s(i) \rightarrow 1$
- If $s(i)$ is small, the point is “between clusters” (or in its own cluster): the distance to the other members of the same cluster is about the same as to other clusters
- If $s(i)$ is negative ($a(i) > b(i)$), the point is seriously misclassified: it’s closer to another cluster than to other members of the cluster it was assigned to









SILHOUETTE: R CODE

- R has a function for calculating the silhouette (in package `cluster`)
 - Watch for average silhouette (cf. all average silhouettes > 0.6 at K=3 (green) and two much weaker clusters at K=4 (blue))

```
kf=kmeans(cbind(x,y),3)
d=dist(cbind(x,y))
sl=silhouette(kf$cluster,d)
summary(sl)
Silhouette of 120 units in 3 clusters ...
Cluster sizes and average silhouette widths:
      40      51      29
0.6316847 0.6408889 0.6410914
Individual silhouette widths:
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
0.1327  0.5657  0.6812  0.6379  0.7287  0.7635
plot(sl,col="lightgreen")
plot(silhouette(kmeans(cbind(x,y),4)$cluster,d),
     col='lightblue')
```



SILHOUETTES OF RANDOM DATA






CODE FOR THE PREVIOUS PLOTS

```
old.par = par(mfrow=c(2,3),ps=16)
d=dist(cbind(x,y))
plot(silhouette(kmeans(cbind(x,y),3,nstart=10)$cluster,d),col='lightblue')
for ( iRnd in 1:5 ) {
  xyrnd = apply(cbind(x,y),2,function(x) runif(length(x),min(x),max(x)))
  drnd = dist(xyrnd)
  plot(silhouette(kmeans(xyrnd,3,nstart=10)$cluster,drnd),col='red')
}
par(old.par)
```

REAL DATA EXAMPLE

- Let us consider again the NCI60 dataset first introduced last week
 - 64 tumor samples, expression levels of 6380 genes are measured in each sample

```
> library(ISLR)
> data(NCI60)
> rownames(NCI60$data)=NCI60$labs
> data.1 = NCI60$data[,apply(NCI60$data,2,sd)>1]
> data.1.s = scale(data.1)
# try k-means first:
> w=numeric(20)
> for ( k in 2:20 ) {
  kf=kmeans(data.1,k)
  w[k] = (kf$betweenss/      # CH index
    (k-1)) / (kf$tot.withinss / (length(kf$cluster)-k))
}
> plot(2:20,w[-1],type="b",lwd=2,pch=19,xlab="K",
  ylab="CH index")
```



! Gap statistics similarly unconvincing!

CLUSTER STABILITY

- How stable are the clusters as we wiggle parameters/rescale the data etc?
- Build a “contingency table” of one clustering result against another
 - Caveat: cluster IDs will be different, need to rearrange the table
 - Problem: given matrix A , find permutation $p: [1 \dots n] \rightarrow [1 \dots n]$ such that the trace of the rearranged matrix, $\sum_i A[i, p(i)]$, reaches its maximum
 - “Hungarian algorithm” provides the solution, but it is not a simple one... But there is a library!

```
matrix.sort <- function(m) {
  require(clue)
  p = solve_LSAP(m, maximum=T) # find the permutation...
  m[, p] # and apply it!
}
```

7	1	2	→	7	2	1
2	0	12		2	12	0
0	9	3		0	3	9

```
cmp.shortcut = function(K) {
  matrix.sort(table(
    FULL=kmeans(NCI60$data, K, 10)$cluster,
    SCALED.SUBSET=
      kmeans(data.1.s, K, 10)$cluster))
}
```

```
cmp.shortcut(4)
SCALED.SUBSET
FULL  1  2  3  4
1    8  0  0  0
4    0 28  0  0
3    0  0  9  0
2    0  2  0 17
```

```
cmp.shortcut(5)
SCALED.SUBSET
FULL  1  2  3  4  5
2    8  0  0  0  1
3    0  9  0  0  0
1    0  0 15  0  0
5    0  0  0  8  0
4    0  0  0  0 23
```

IS HIERARCHICAL BETTER?

- Functions for within/between cluster variance:




```
within=function(d,clust) {  
  w=numeric(length(unique(clust)))  
  for ( i in sort(unique(clust)) ) {  
    members = d[clust==i,,drop=F]  
    centroid = apply(members,2,mean)  
    members.diff = sweep(members,2,centroid)  
    w[i] = sum(members.diff^2)  
  }  
  return(w)  
}
```

```
between=function(d,clust) {  
  b=0  
  total.mean = apply(d,2,mean)  
  for ( i in sort(unique(clust)) ) {  
    members = d[clust==i,,drop=F]  
    centroid = apply(members,2,mean)  
    b = b + nrow(members)*  
        sum( (centroid-total.mean)^2 )  
  }  
  return(b)  
}
```

EXAMINING DIFFERENT NUMBERS OF HIERARCHICAL CLUSTERS

- The clear cluster structure is missing from the hierarchical clustering result as well...


```
data.1=NCI60$data[,apply(NCI60$data,2,sd)>1]
dd.1=dist(data.1)
hw.1=hclust(dd.1,method="ward.D2")
w.tot=numeric(9)
btw=numeric(9)
for ( k in 2:10 ) {
  clust = cutree(hw.1,k=k)
  w = within(data.1,clust)
  w.tot[k-1]=sum(w)
  btw[k-1] = between(data.1,clust)
}
plot(2:10,w.tot,pch=19,type="b")
plot(2:10,btw,pch=19,type="b")
plot(2:10,(btw/(1:9))/
      (w.tot/(nrow(data.1)-2:10)),pch=19,type="b")
```



BRUTE FORCE RANDOMIZATION


- Let us randomize expression levels across each sample and repeat clustering
- All the “similarities” between samples we are now observing are due to random chance
- What is the distribution of the distances between samples/clusters and how does it compare to the original data?

```
ori.heights = hw.1$height
rnd.heights = numeric()
for ( i.sim in 1:100 ) {
  data.rnd <- apply(data.1, 2, sample)
  hw.rnd = hclust(dist(data.rnd), method="ward.D2")
  rnd.heights <- c(rnd.heights, hw.rnd$height)
}
plot(ori.heights, rank(ori.heights)/length(ori.heights),
     col="red", xlab="height", ylab="F(height)", pch=19)
points(rnd.heights, rank(rnd.heights)/length(rnd.heights),
       col="blue")
```




NCI60 CLUSTERS

- About 4 top level clusters *might* be real
 - Brute force resampling can be too optimistic!
- Many other methods exist, including bootstrap of variables (genes in our example) or multilevel bootstrap
 - See package `pvclust`




NOTE ON DISTANCES




Distance similarity measure

(a) Correlation coefficient without centering



(c) Absolute correlation coefficient without centering



(d) Absolute correlation coefficient with centering



(e) Euclidean distance



(f) Manhattan distance



TRANSFORM in GeneRank

“AIRWAY” DATASET

- Another gene expression experiment
 - Can be downloaded as library/dataset from Bioconductor (another large repository of R packages, biology-oriented)
 - 4 cell lines (lung airway, asthma), 2 conditions: untreated (control) and treated with a drug


```
source("http://bioconductor.org/biocLite.R")
biocLite("airway")
library(airway)
data(airway)
airway
class: SummarizedExperiment
dim: 64102 8
exptData(1): ''
assays(1): counts
rownames(64102): ENSG000000000003 ENSG000000000005 ...
rowData metadata column names(0):
colnames(8): SRR1039508 SRR1039509 ...
colData names(9): SampleName cell ...
```

```
> colData(airway)[1:3,1:4]
DataFrame with 3 rows and 4 columns
      SampleName      cell      dex      albut
      <factor> <factor> <factor> <factor>
SRR1039508 GSM1275862   N61311   untrt   untrt
SRR1039509 GSM1275863   N61311    trt    untrt
SRR1039512 GSM1275866  N052611   untrt   untrt
> assay(airway)[1:3,1:3]
      SRR1039508 SRR1039509 SRR1039512
ENSG000000000003      679      448      873
ENSG000000000005        0        0        0
ENSG000000000419      467      515      621
```

PREPARE DATA

- Many genes are in fact not measured (count = 0 or extremely low) : remove unmeasured genes
- If we “count longer”, we will count proportionally more events per gene. Need to normalize the counts by the total number of counts in each sample
 - Better methods do exist!
- Distribution of expression levels is log-normal. Take the log.


```
> m = assay(airway)
# use cell line/treatment to rename columns (samples):
> colnames(m) = paste(colData(airway)$cell,
                      ifelse(colData(airway)$dex=="untrt", "U", "T"), sep=".")
> labels=colnames(m)
> m = m[ apply(m,1,sum) > 5, ] # cuts m from 64K rows (genes) to 24K
> total.counts=apply(m,2,sum)
> m = sweep(m,2,total.counts/1000000,FUN="/") # normalize per MILLION cnts
> m=log2(m+1) # regularize log
> boxplot(m,col="lightblue",las=3)
```



CLUSTER SAMPLES

- Let us consider the genes (“variables”) with largest variance across the “experimental observations” (samples) and cluster samples on those genes
 - Not shown: we could cluster on full set of variables, the results would be less stable but still strange



```
# correlation distance (Note the t(): we are going to use the
# same contract as dist() which calculates pairwise
# distances between rows (observations)
corr.dist=function(x) { as.dist(1-cor(t(x))) }
vars=apply(m,1,var)
slct=order(vars,decreasing=T)[1:300]
plot(hclust(corr.dist(t(m[slct,])),method="ward.D2"),
     main="Correlation, Ward, top 300 vars")
plot(hclust(dist(t(m[slct,])),method="average"),
     main="Euclidean, Average, top 300 vars")
```



INTRODUCING HEATMAP

- Simplest heatmap
 - Order of rows/columns the same as in the data table (must use Rowv=F and Colv=F); values are represented with color gradient (see the legend)
 - Note how we set color code for untreated/treated (green/red) on top (ColSideColors)


```
library(gplots)
slct=order(vars,decreasing=T)[1:20]
heatmap.2(m[slct,],trace="none",labCol=labels,
          ColSideColors=rep(c("green","red"),4),
          margins=c(7,7),Rowv=F,Colv=F)
```



CLUSTERED HEATMAP

- Heatmap can also automatically cluster both rows AND columns
- Two-way clustering: samples are clustered according to the similarity of expression patterns across (considered) genes in each sample
- Genes are clustered according to the similarity of their expression patterns across the samples

```
heatmap.2(m[slect,], trace="none",  
  labCol=labels,  
  ColSideColors=rep(c("green", "red"), 4),  
  margins=c(7, 7))
```



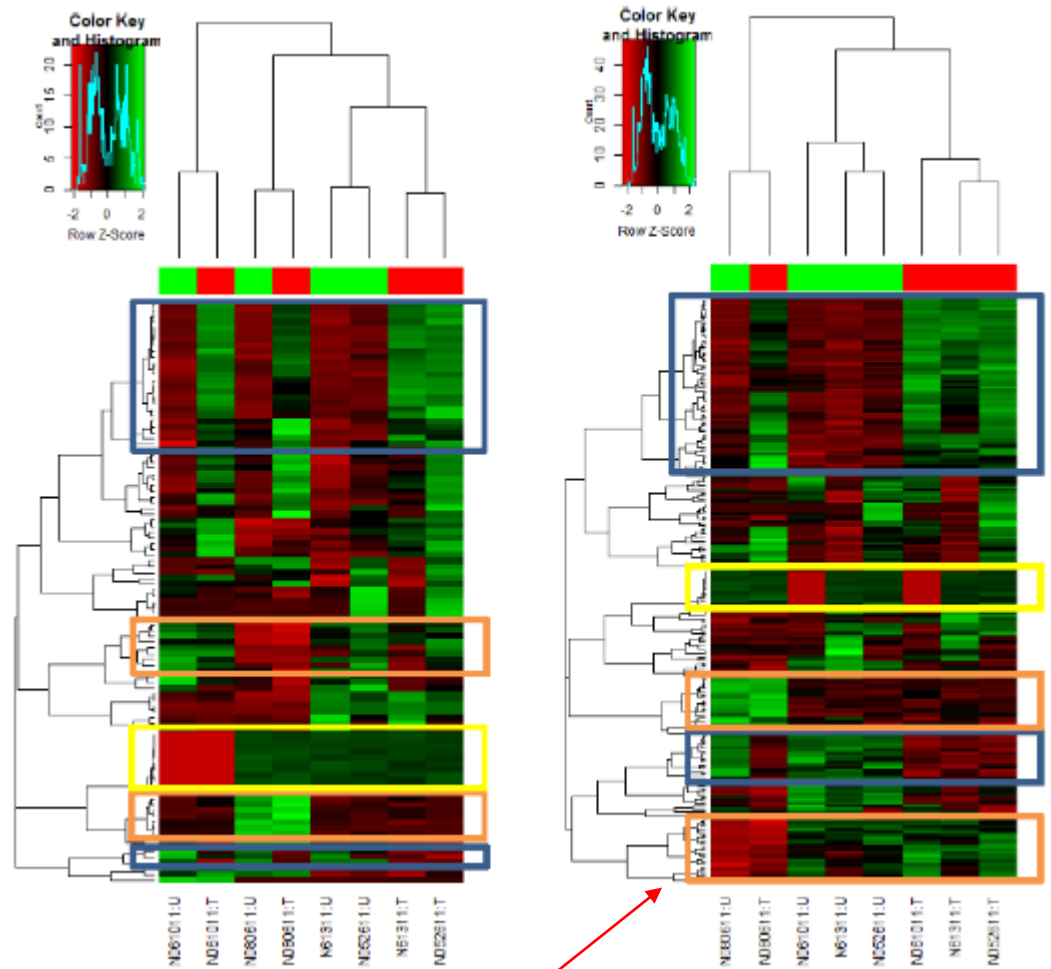
CONSIDERING LARGER SET

- Let us (a) consider the larger set of genes (comparable to what we used before), (b) use the same clustering distance (heatmap allows us to do that!) , and (c) enhance visualization:

```
# generates a vector of n colors transitioning from red, through black,  
# to green  
redgreen <- function(n) {  
  c( hsv(h=0/6, v=seq(1,0,length=n/2) ),hsv(h=2/6, v=seq(0,1,length=n/2) ) )  
}  
# try 100 highest-variance genes first  
slct=order(vars,decreasing=T)[1:100]  
# use correlation distance, center and scale rows, use red-green palette,  
# and turn off gene labels for now - there are too many to be readable anyway:  
heatmap.2(m[slct,],trace="none",labCol=labels, labRow=F,  
          col=redgreen(100),scale="row",  
          ColSideColors=rep(c("green","red"),4),margins=c(7,7),distfun=corr.dist)  
#now try 200 highest-variance genes  
slct=order(vars,decreasing=T)[1:200]  
heatmap.2(m[slct,],trace="none",labCol=labels,labRow=F,  
          col=redgreen(100),scale="row",  
          ColSideColors=rep(c("green","red"),4),margins=c(7,7),distfun=corr.dist)
```

DISTINCT PATTERS OF EXPRESSION

- We observe the same effect: one cell line starts clustering away when we consider 200 genes
- We clearly see a few different gene clusters with specific expression patterns.
 - Some distinct clusters are outlined with colored rectangles, note different expression patterns across samples
 - Note the cluster that differentiates the outlier sample from the rest
 - Upon close examination, we would discover that downregulated genes in this cluster are located... on Y-chromosome!
 - In contrast to *what's claimed in the publication*, one cell line comes from a female!



SUMMARY

- The question “how many clusters” is *very difficult*
 - Just like everything else we have seen so far: the question is not only “how many clusters are truly there” but also “how many clusters we can possibly distinguish with the data in hand” ! (cf.: “what’s the true underlying dependence” vs “what we can possibly do with available data” – for instance should we even try higher order polynomial regression, or that would give us too many variables for the small dataset we have?)
 - There is a lot of metrics and statistics that can (and should) be examined (again, similar to the large number of statistics we may want to look at when performing model selection)
- The question of “how good the clusters are” is closely related (when the “goodness” metrics are clear and telling, we are also likely to know how many clusters are there!)
- Nonetheless, unsupervised learning (e.g. clustering) is very important and useful for understanding the structure of the data, sample QC, and for discovery/hypothesis generation
- Another potential use: a *recommender system* (e.g. find customers who are “similar” to customer X and send X offers that other members of their cluster were responding to)